

Integrity Checking

Translating Integrity Rules to SQL

Chuan LU
(With credit to Nigel Hardy)

Modelling Persistent Data
CS27020
Computer Science Department

Outline

1 Domain Constraints

2 Referential Integrity - Simple Cases

Domains

- The relational model says that domains must exist
- SQL specifies a minimum list of data types
 - Varies with versions of SQL
 - https://en.wikibooks.org/wiki/Structured_Query_Language/Data_Types
- Each DBMS will offer these plus possibly more
 - Postgres: <https://www.postgresql.org/docs/11/datatype.html>
- Some attributes may benefit from tighter constraint on values

Attributes with constraints

- Limiting acceptable values when the table is created
 - or “ALTER”ed
- (NOT NULL)
- CHECK clauses can be added
 - to individual attributes
 - to the table, concerning attributes

Simple CHECK

```
CREATE TABLE products (  
    product_no integer ,  
    name text ,  
    price numeric CHECK (price > 0)  
);
```

Credit: Postgresql documentation

<https://www.postgresql.org/docs/11/ddl-constraints.html>

Named constraint

```
CREATE TABLE products (  
    product_no integer ,  
    name text ,  
    price numeric  
    CONSTRAINT positive_price CHECK (price > 0)  
);
```

Credit: Postgresql documentation

<https://www.postgresql.org/docs/11/ddl-constraints.html>

Table constraints

```
CREATE TABLE products (  
    product_no integer,  
    name text,  
    price numeric CHECK (price > 0),  
    discounted_price numeric CHECK (discounted_price > 0),  
    CHECK (price > discounted_price)  
);
```

Credit: Postgresql documentation

<https://www.postgresql.org/docs/9.6/ddl-constraints.html>

Application program constraint

- Domain constraint could be left to the application program
- Duplicate code, error prone
- Possible that constraints vary with time or circumstance
- Double check may be beneficial
 - “front end” feedback

Creating new domains

- Multiple use

```
CREATE DOMAIN name [ AS ] data_type
    [ COLLATE collation ]
    [ DEFAULT expression ]
    [ constraint [ ... ] ]

where constraint is:

[ CONSTRAINT constraint_name ]
{ NOT NULL | NULL | CHECK (expression) }
```

Credit: PostgreSQL documentation

<https://www.postgresql.org/docs/9.6/ddl-constraints.html>

New domain example

```
DROP DOMAIN IF EXISTS cooktemp;

CREATE DOMAIN cooktemp AS INT
  CHECK ( VALUE >= 100 AND VALUE <=250);

CREATE TABLE receipe (
  name VARCHAR(30) PRIMARY KEY,
  temp cooktemp
  — and obviously other attributes
);

INSERT INTO receipe VALUES ( 'Jam Tarts' ,180);
INSERT INTO receipe VALUES ( 'Mince Pies' ,1800);
```

IN construct

- “enumeration”

```
CREATE TABLE students (  
  id VARCHAR(9) PRIMARY KEY,  
  name VARCHAR(30),  
  — and other attributes  
  level VARCHAR(3) CHECK IN ( 'UG' , 'PGT' , 'PGR' )  
);
```

- Is the list fixed permanently?
- Would an FK be better?

Outline

1 Domain Constraints


2 Referential Integrity - Simple Cases

Outline

1 Domain Constraints

2 Referential Integrity - Simple Cases

One to Many (Optional Many)



Integrity Rules: A module must be run by a single department; however, a department can have zero to many modules (i.e. service departments vs. academic departments).


The “many” end contains a foreign key reference to the “one” end.

One to Many (Optional Many) - SQL implementation

```
CREATE TABLE Department(  
    id VARCHAR(5) PRIMARY KEY,  
    name VARCHAR(50),  
    UNIQUE(id, name));  
  
CREATE TABLE Module(  
    id VARCHAR(7) PRIMARY KEY,  
    deptID VARCHAR(5) NOT NULL  
        REFERENCES Department(id),  
    title VARCHAR(50),  
    UNIQUE(id, deptID, title));
```

(70_DeptModule.sql)


One to Many (Optional Many)



Cardinality implementations:

0	There need be no module record with a particular department PK planted
*	There can be many module records with a particular department PK planted
1	ID is a PK - it cannot be null
1	ID is a PK - it cannot be duplicated

One to Many (Optional Many)



Cardinality implementations:

0	There need be no module record with a particular department PK planted
*	There can be many module records with a particular department PK planted
1	ID is a PK - it cannot be null
1	ID is a PK - it cannot be duplicated

(In all these examples for “PK” read “PK or CK”)

Many to Many (Optional Both Sides)



Integrity Rules: A student may take no modules, or may take one or more modules. A module may have no registered students, or may have several.

The many-to-many relationship becomes a new relation with a composite key (the *takes* relationship becomes the Takes relation).


(StudentModule.sql)

Many to Many (Optional Both Sides) - SQL implementation

```
CREATE TABLE Student(  
    studentID VARCHAR(5) PRIMARY KEY,  
    firstname VARCHAR(50),  
    surname VARCHAR(50));  
  
CREATE TABLE Module(  
    id VARCHAR(7) PRIMARY KEY,  
    title VARCHAR(50),  
    UNIQUE(id, title));  
  
CREATE TABLE Takes (  
    studID VARCHAR(5)  
        REFERENCES Student(studentID),  
    modID VARCHAR(7)  
        REFERENCES Module(id),  
    PRIMARY KEY (studID, modID)  
);
```

(80_StudentModule.sql)

Many to Many (Optional Both Sides)




Cardinality implementations:

0	There need be no takes record with a particular module PK planted
*	There can be many takes records with a particular module PK planted
0	There need be no takes record with a particular student PK planted
*	There can be many takes records with a particular student PK planted

student cannot take the same module twice: PK of takes

One to N (Optional Both Sides)



Integrity Rules: A student may have at most one sponsor. A sponsor may have no students, or may have several.

Create a new relation with same key as 'many' end.

(StudentSponsor.sql)

One to N (Optional Both Sides) - SQL implementation


```
CREATE TABLE Student(  
    studentID VARCHAR(5) PRIMARY KEY,  
    firstname VARCHAR(50),  
    surname VARCHAR(50));
```

```
CREATE TABLE Sponsor(  
    companyID VARCHAR(6) PRIMARY KEY,  
    companyName TEXT,  
    UNIQUE(companyID, companyName));
```

```
CREATE TABLE Sponsors(  
    studentID VARCHAR(5)  
        REFERENCES Student(studentID),  
    companyID VARCHAR(6)  
        REFERENCES Sponsor(companyID),  
    PRIMARY KEY (studentID));
```

(90_StudentSponsor.sql)

One to N (Optional Both Sides)



Cardinality implementations:

0	There need be no sponsors record with a particular student PK planted
1	There can only be 1 sponsors record with a particular student PK planted - PK (unique)
0	There need be no sponsors record with a particular sponsor PK planted
*	There can be many sponsors records with a particular sponsor PK planted

Alternative One to N (Optional Both Sides)

Allow a NULL FK

Alternative One to N (Optional Both Sides)

Allow a NULL FK

```
CREATE TABLE Sponsor(  
  companyID VARCHAR(6) PRIMARY KEY,  
  companyName TEXT,  
  UNIQUE(companyID, companyName));
```

```
CREATE TABLE Student(  
  studentID VARCHAR(5) PRIMARY KEY,  
  firstname VARCHAR(50),  
  surname VARCHAR(50));
```

```
CREATE TABLE Sponsors(  
  studentID VARCHAR(5)  
    REFERENCES Student(studentID),  
  companyID VARCHAR(6) NOT NULL  
    REFERENCES Sponsor(companyID),  
  PRIMARY KEY (studentID));
```

Alternative One to N (Optional Both Sides)

Allow a NULL FK

```
CREATE TABLE Sponsor(  
  companyID VARCHAR(6) PRIMARY KEY,  
  companyName TEXT,  
  UNIQUE(companyID, companyName));
```

```
CREATE TABLE Student(  
  studentID VARCHAR(5) PRIMARY KEY,  
  firstname VARCHAR(50),  
  surname VARCHAR(50));
```

```
CREATE TABLE Sponsors(  
  studentID VARCHAR(5)  
    REFERENCES Student(studentID),  
  companyID VARCHAR(6) NOT NULL  
    REFERENCES Sponsor(companyID),  
  PRIMARY KEY (studentID));
```

```
CREATE TABLE Sponsor(  
  companyID VARCHAR(6) PRIMARY KEY,  
  companyName TEXT,  
  UNIQUE(companyID, companyName));
```

```
CREATE TABLE Student(  
  studentID VARCHAR(5) PRIMARY KEY,  
  firstname VARCHAR(50),  
  surname VARCHAR(50)  
  sponsor VARCHAR(6)  
    REFERENCES Sponsor(companyID);
```

Integrity Checking so far

- UNIQUE
- NOT NULL
- Build-in domains (types)
- CHECK
- new DOMAIN
- Referential Integrity for
 - 1:M (Optional many)
 - M:N (Optional on Both Sides)
 - 1:N (Optional on Both Sides)

Code available from Blackboard